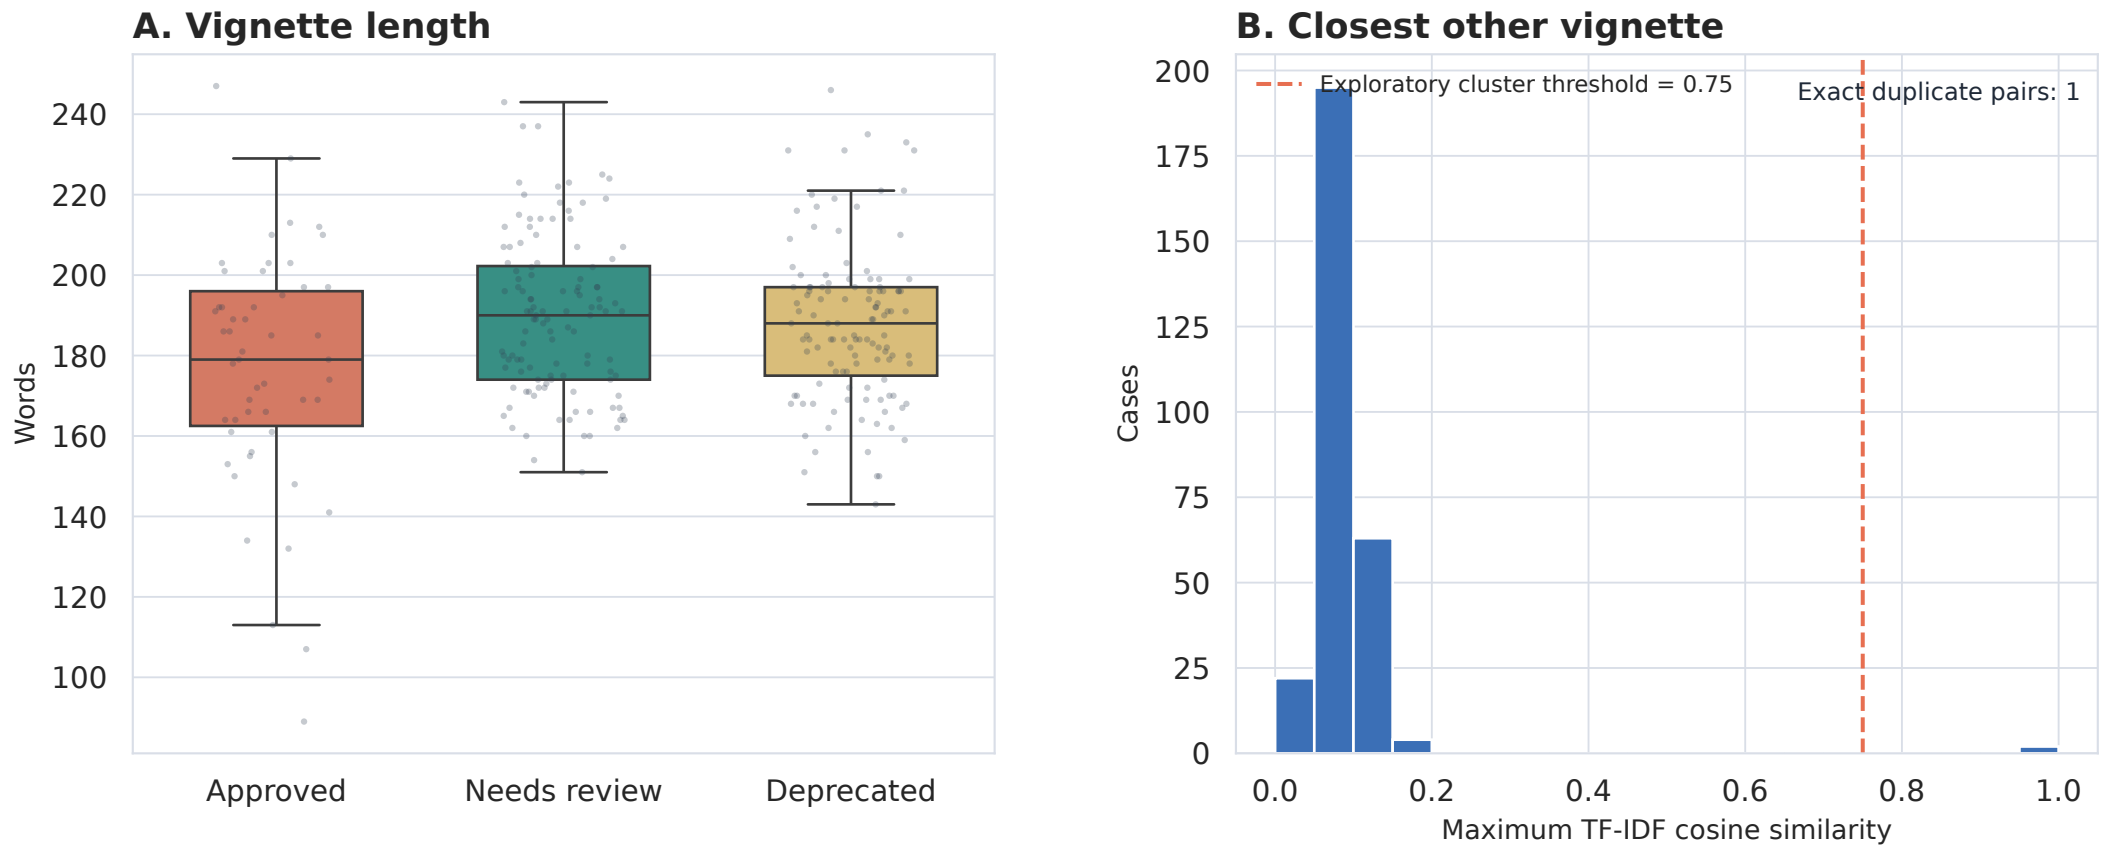


Figure 9. Text audit identifies length variation and potential split leakage



TF-IDF similarity is a screening tool. High-similarity pairs require manual review before assigning group-aware train/validation splits.

Source: ValueBench EDA Google Sheet (<https://docs.google.com/spreadsheets/d/1-Zhf4Nr59UqoeJLqhFwLoGDRH8uAuMGKdNqnigVO01Y/edit?gid=220762358>).